

B+TREE

QuyênPT3-TranNDC



Agenda

B-TREE

B-Tree structure
Pros and cons

Implement Key Value
Store with B+TREE
- Files

01

02

03

04

B+TREE

B+Tree structure
Why?
Applications

DEMO



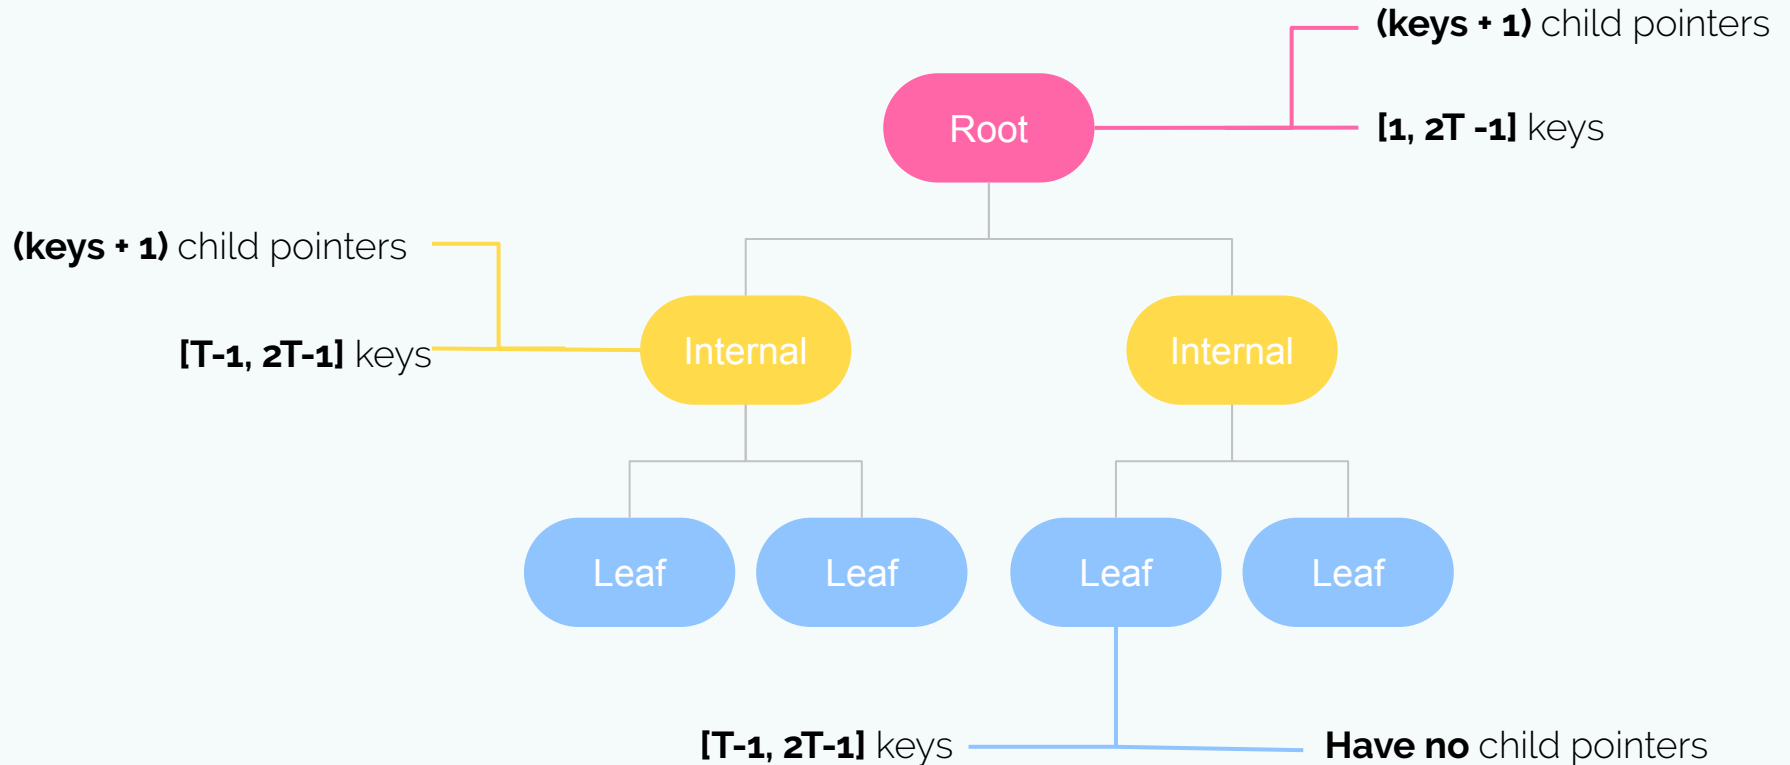


01

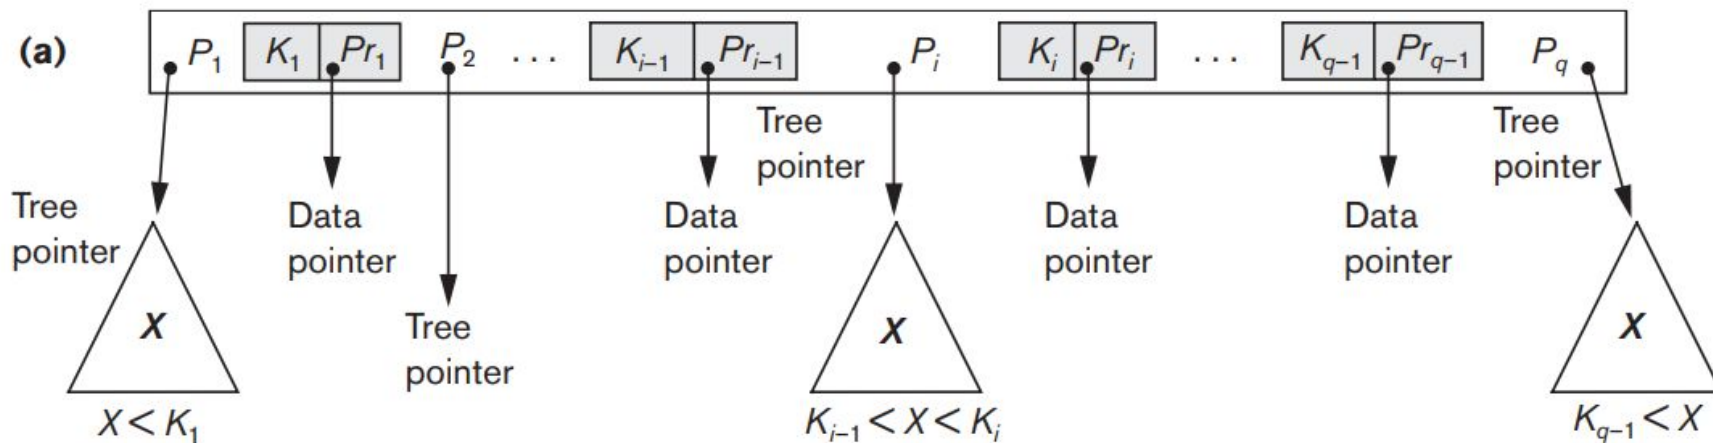
B-TREE

B-Tree Structure
Pros and cons

B-Tree Structure



B-Tree Structure



Pros and Cons



- **Height of the tree** is minimized
- Reduce disk latency: **$O(h)$ disk access**
- Total CPU time used: **$O(th)$**



- The **number of keys** in one node is **limited** because keys are stored along with their values (or pointers of values)
- Poor write performance: Random writes and write-amplification
- Delete operation is complex when handling **deletion on internal nodes**

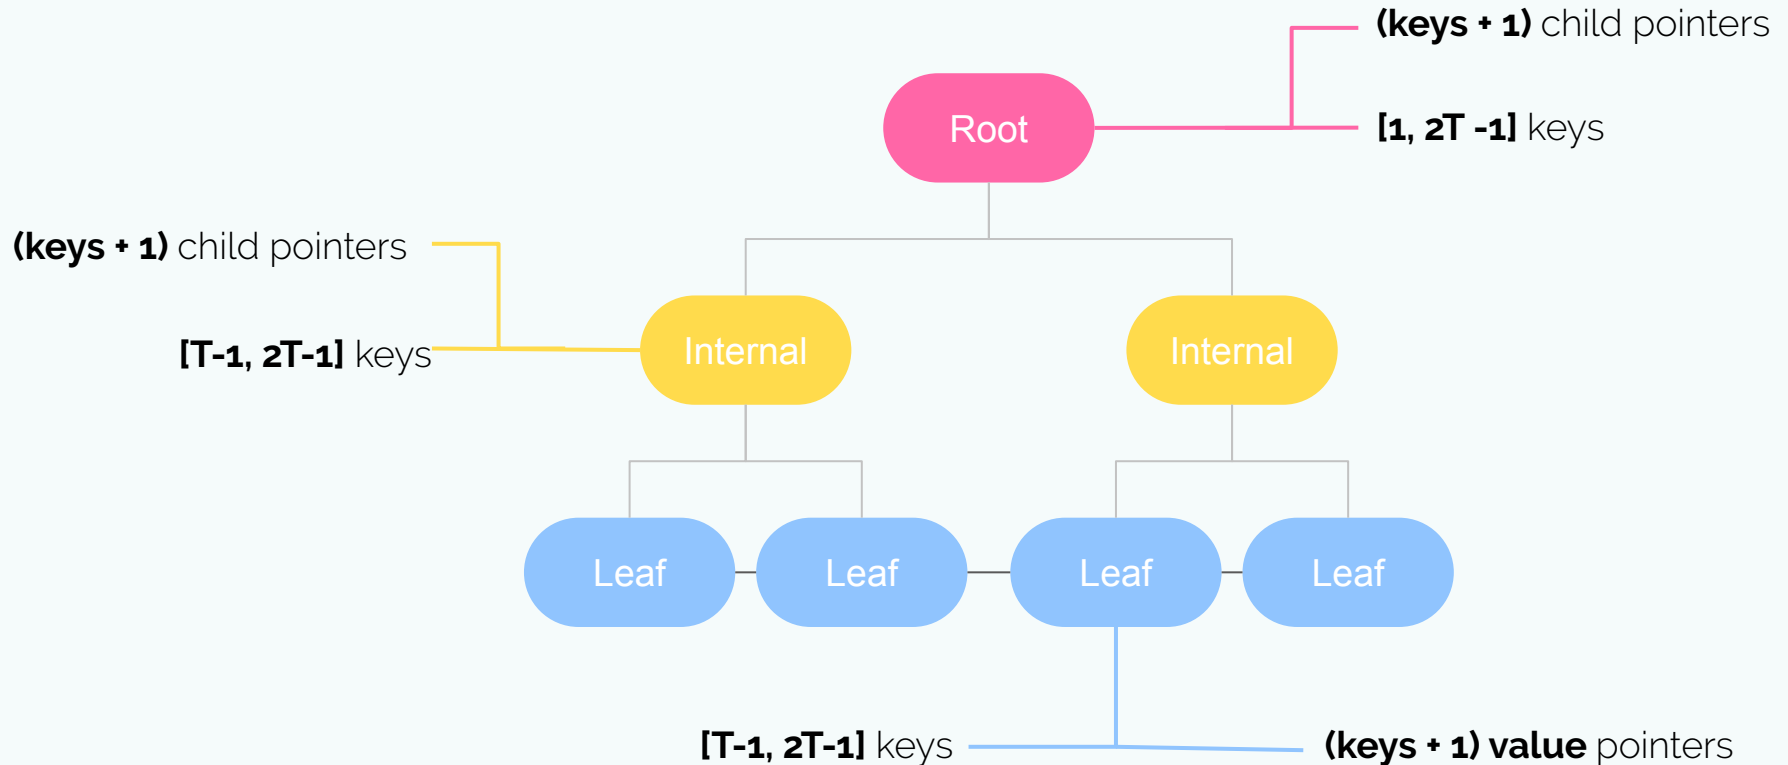
02

B+TREE

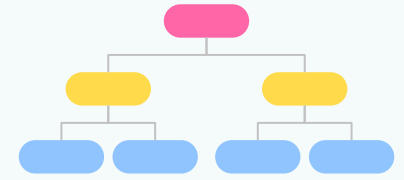
B+Tree Structure
Why?
Applications



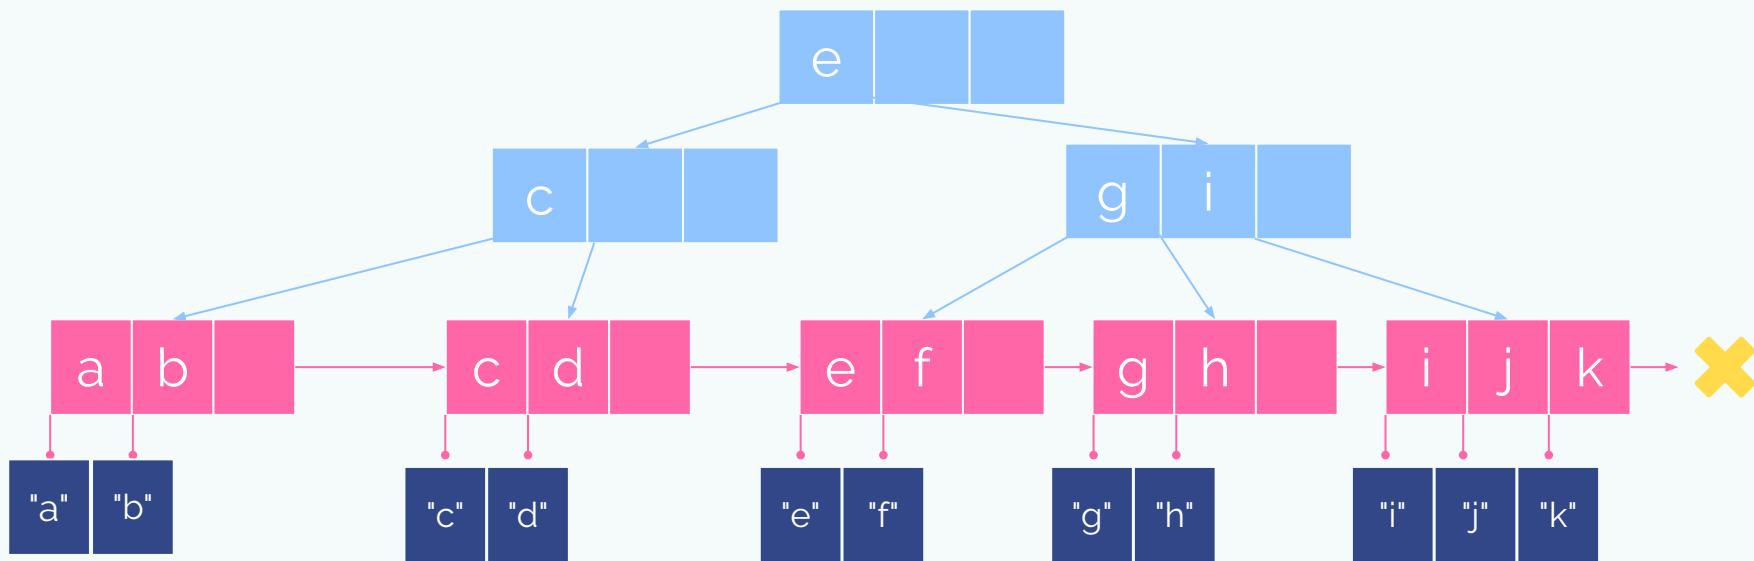
B+Tree Structure



B+Tree Structure




B+Tree Example



B+Tree operations complexity

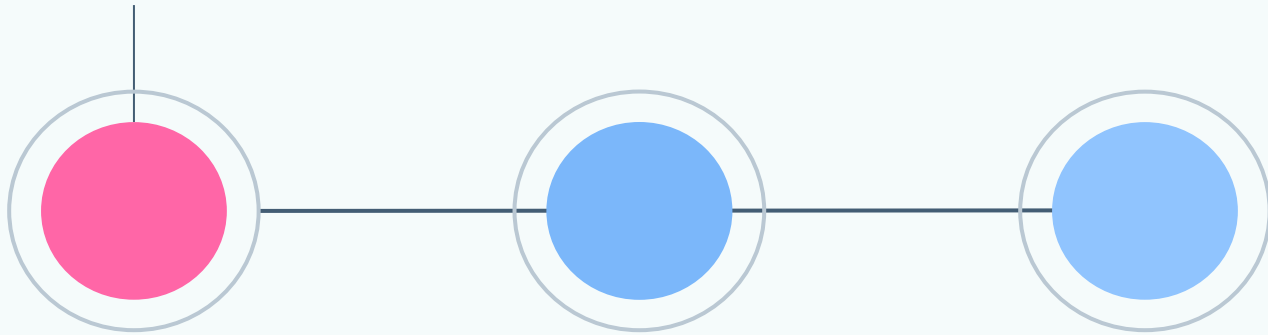
- **Insert:** $O(\log N)$
- **Remove:** $O(\log N)$
- **Get:** $O(\log N)$
- **Exist:** $O(\log N)$



Why B+Tree ?

Why B+Tree ?

Structure of nodes
are **unified**

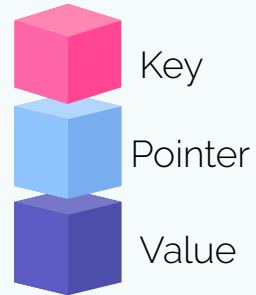


B-Tree

Internal Node



Leaf Node

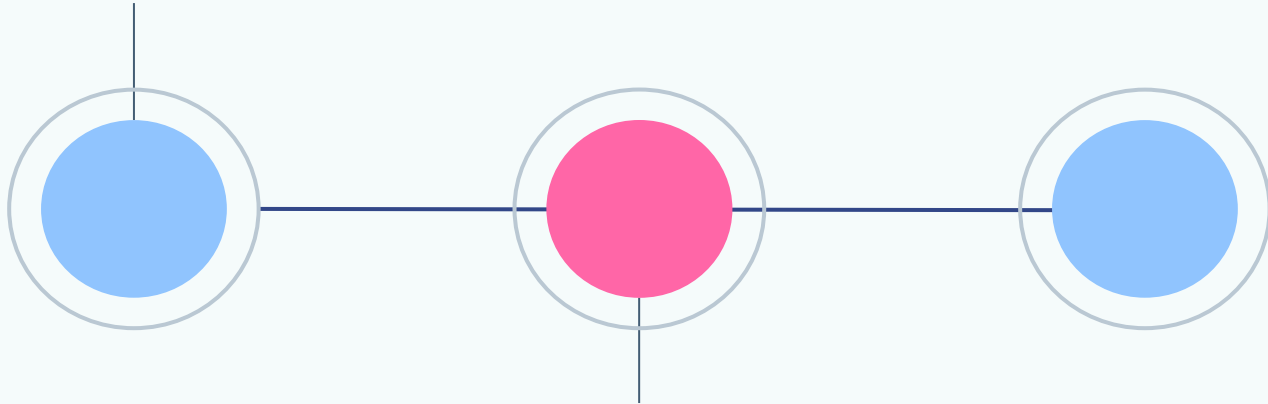


B+Tree



Why B+Tree ?

Structure of nodes
are **unified**



Increase branching factor:
reduce height of tree
comparing with B-Tree

- **Key's size:** 4 bytes
- **Value pointers:** 8 bytes
- **Node pointers:** 8 bytes

B-TREE

	NUM NODES	NUM KEYS	NUM VALUE POINTERS	NUM NODE POINTERS
ROOT	1	50	50	51
LEVEL 1	51	2550	2550	2601
LEVEL 2	132651	130050	130050	-



- **Key's size:** 4 bytes
- **Value pointers:** 8 bytes
- **Node pointers:** 8 bytes

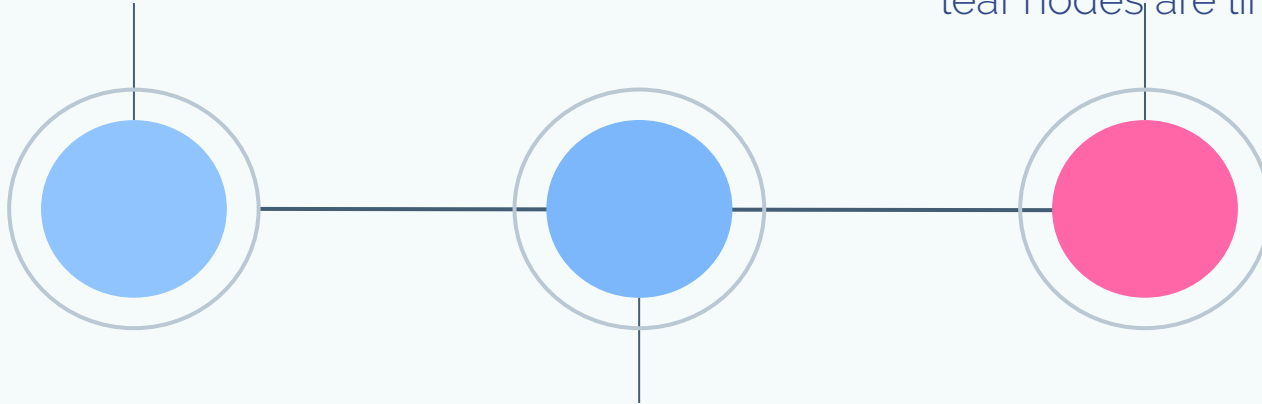
B+TREE

	NUM NODES	NUM KEYS	NUM NODE POINTERS
ROOT	1	84	85
LEVEL 1	85	7140	7225
LEVEL 2	7225	599675	614125



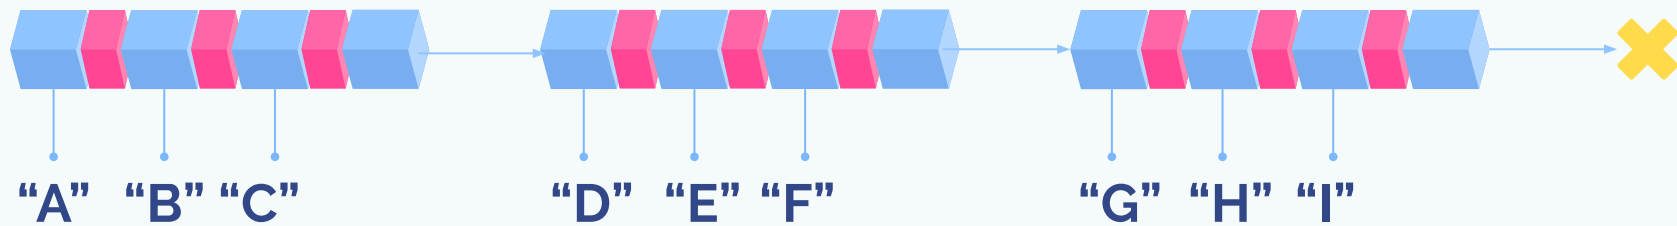
Why B+Tree ?

Structure of nodes
are **unified**



Increase branching factor:
reduce height of tree
comparing with B-Tree

Increase **range queries performance** because
leaf nodes are linked



B+Tree's applications





Implement Key-Value Store with B+TREE

03

FILES



DATABASE

Store Key-Value entries



DELETED NODES

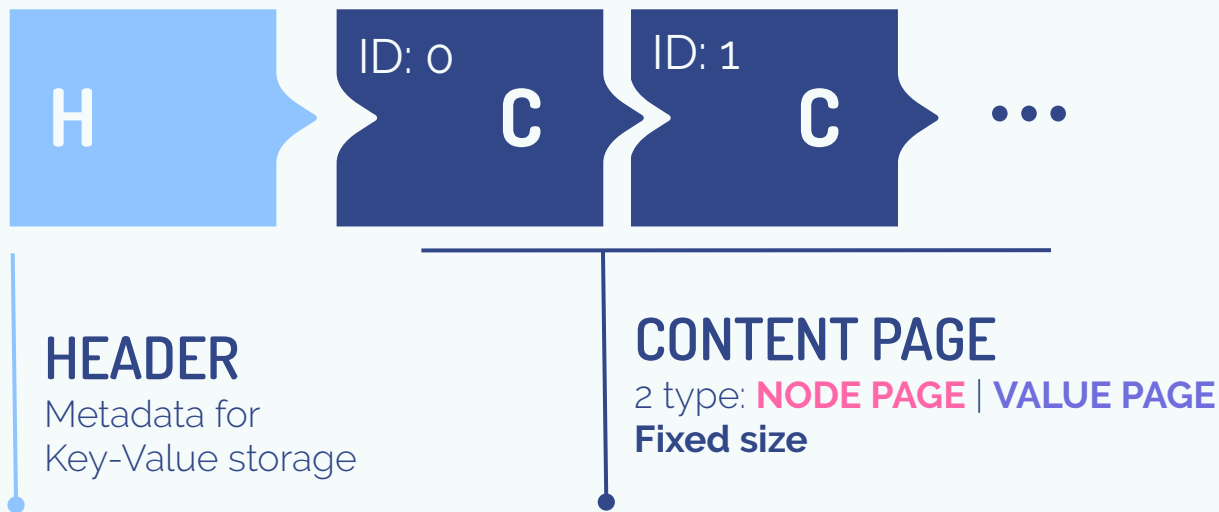
Offsets of deleted nodes



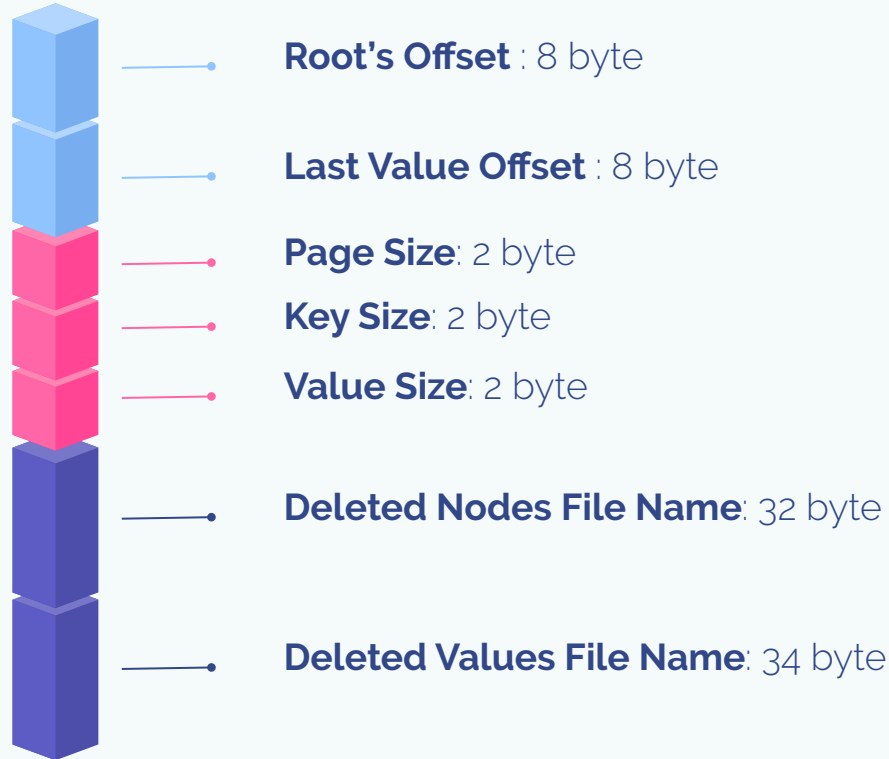
DELETED VALUES

Offsets of deleted values

DATABASE FILE STRUCTURE

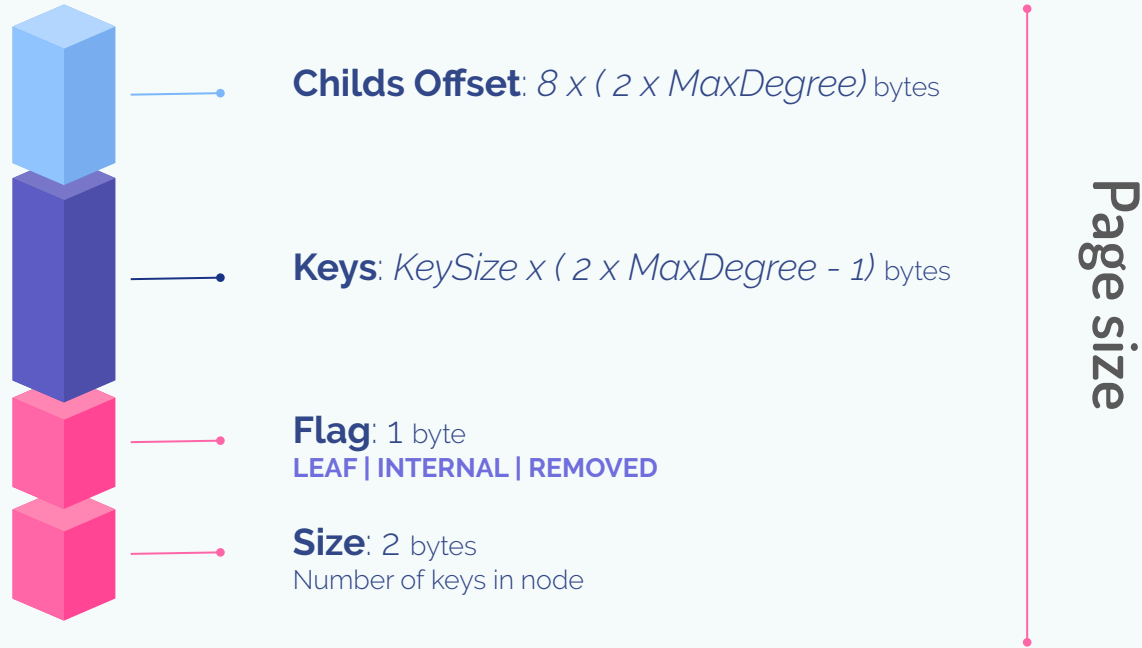


Header

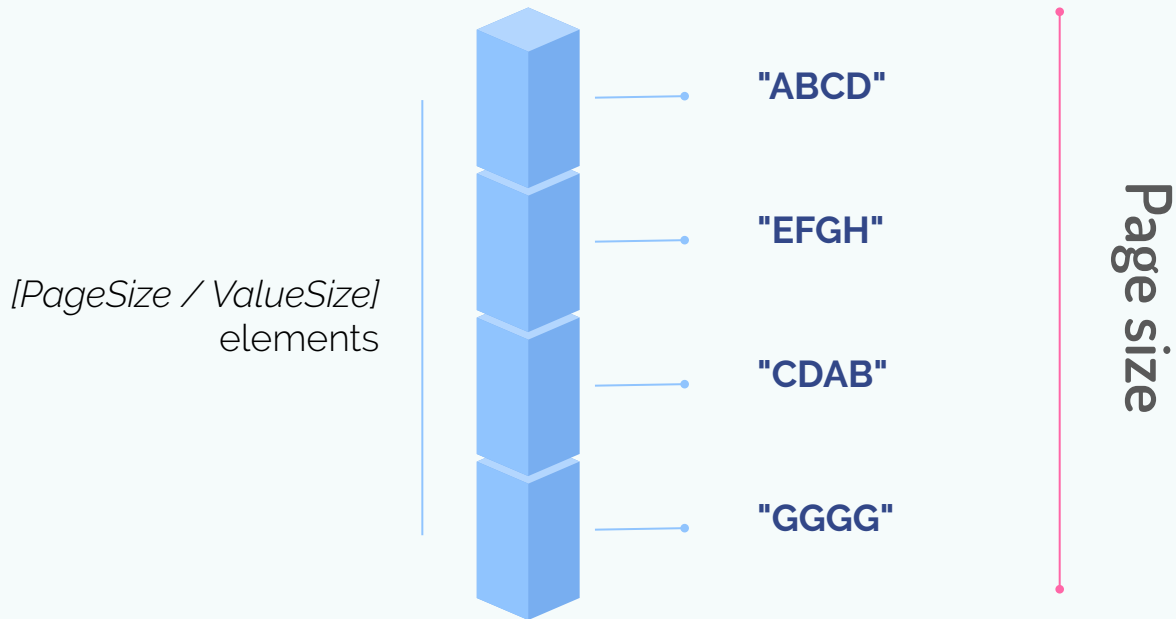


88 bytes

Node page



Value page



Offset

INDEX OF VALUE IN NODE

$\text{bitSize}(\text{PageSize} / \text{ValueSize})$ bits
= 00 If it's NodePage

REAL NODE ID

$64 - \text{bitSize}(\text{PageSize} / \text{ValueSize})$ bits

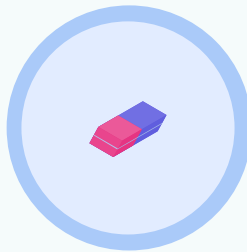


FILES



DATABASE

Store Key-Value entries



DELETED NODES

Offsets of deleted nodes



DELETED VALUES

Offsets of deleted values

DELETED FILE STRUCTURE





DEMO

04



Q/A

THANKS FOR
LISTENING

